
DPGrammar: Joint Viterbi Repair for Grammar-Constrained Diffusion Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Constrained decoding keeps LLM output in a format like JSON by masking invalid
2 tokens with a left-to-right parser. Diffusion language models unmask by confidence
3 rather than in order, so a token is final before the parser reaches it: the grammar
4 cannot be enforced ahead of the parse frontier, only repaired behind it. One position
5 at a time is not enough, because a grammar couples the positions it constrains
6 and a locally probable fix can eliminate every valid completion of the region.
7 **DPGrammar** decides the violated span as a whole. Merging paths that leave
8 the parser in the same state, it runs a Viterbi pass over (parser state $\times$ position)
9 pairs and returns the likeliest valid assignment in the lattice it expands, in time
10 linear rather than exponential in the span length, with candidates drawn from the
11 automaton’s outgoing edges rather than the model’s ranking, so a tighter budget
12 costs likelihood and never validity. On JSONSchemaBench medium with LLaDA-
13 8B-Instruct, schema validity rises from 91.8% to 96.7% and hand-backs fall from
14 33.2 to 1.2 per instance, for 3.3 s of added mean latency spent writing $1.7\times$ the
15 tokens where the layer fires. The margin over the unrepaired baseline more than
16 doubles, to 10.3 pp, once outputs must carry content, and reaches 13.4 pp on the
17 hardest split.

18 1 Introduction

19 Masked diffusion language models (dLLMs) such as LLaDA-8B [Nie et al., 2025] and Dream-7B
20 [Ye et al., 2025] generate by iteratively unmasking positions in confidence order rather than left to
21 right, decoding many tokens per forward pass and trading sequential depth for parallelism. Many
22 applications need that output to be structured: an agent calling a tool needs JSON conforming to a
23 given schema [Patil et al., 2025, Geng et al., 2025], and one misplaced brace makes the response
24 unusable to the caller no matter how good its content. Autoregressive (AR) models meet this
25 requirement through constrained decoding: an incremental parser tracks the partial output and masks
26 any next token that would violate the grammar [Willard and Louf, 2023], machinery now standard in
27 serving frameworks [Dong et al., 2024, Guidance AI, 2023, Zheng et al., 2024, Microsoft, 2025].

28 That machinery assumes left-to-right commitment, whereas a dLLM commits out of order: under
29 carry-over unmasking [Sahoo et al., 2024, Shi et al., 2024], once a position is revealed it remains
30 fixed, so the parser may eventually reach an earlier token it rejects but can no longer modify. One
31 answer is to prevent violations outright, making the constraint globally tractable and sampling from
32 the constrained posterior exactly [Suresh et al., 2025]. That guarantees satisfaction by construction,
33 but it needs a regular representation, and a recursive JSON Schema with unbounded nesting is not

regular.¹ It also pays at every denoising step and fixes the decoding order, where repair costs nothing on a clean step and leaves the schedule free. We therefore repair after the fact, forgoing the certified guarantee. Existing repair, though, is myopic [Zhang et al., 2026, Mündler et al., 2025a]: it rewrites the position the parser rejected and leaves the rest of the region alone. That is not enough, because whether a token leads anywhere depends on what must follow it: a choice that is both locally probable and locally legal can still eliminate every valid completion of the region. The model gives no signal about this, its per-position distributions being independent where the grammar is not.

We introduce **DPGrammar**, a repair layer that treats a violated region as a single decision problem rather than a sequence of independent ones. When the parser rejects at the frontier, we identify the affected span and run a Viterbi search over a lattice of (parser state $\times$ position) pairs. Deciding an affected region jointly would be exponential in span length if searched over raw token sequences. However, because a path’s future valid continuations are determined by its current parser state rather than its prefix, paths arriving at identical parser states can merge into a single search node, so the pass costs $O(\text{span} \cdot |S| \cdot k)$ for $|S|$ distinct states per layer, linear in the span rather than exponential in it. We key states on the token set the parser admits, which merges more aggressively than that argument licenses; §3.3 measures what the approximation costs. At each node, candidate tokens are drawn from the parser automaton’s outgoing edges and ranked by model likelihood. Every candidate is one the parser will accept, so a tighter budget costs likelihood and never validity.² Termination is likewise read off the grammar rather than a step count; the layer places no requirement on the unmasking schedule, since it fires only after a rejection; and delegating the parse to an off-the-shelf incremental parser [Microsoft, 2025] makes the constraint context-free rather than regular.

We evaluate on JSONSchemaBench [Geng et al., 2025] and JSON-Mode-Eval [Nous Research, 2024, ETH SRI, 2025] with LLaDA-8B-Instruct, against published constrained decoders and against the same system with the repair layer disabled. On 511 medium instances, DPGrammar reaches 96.7% schema validity against 91.8% with the layer off, and cuts mean repair events from 33.2 to 1.2. The layer’s benefit is not a constant offset but grows with the density of joint constraints, from +0.2 to +13.4 pp of validity between the easiest and hardest split. On a grammar whose violations are purely local the layer never fires at all, which places the same claim at the other end of that axis.

2 Related Work

Constrained decoding for AR LLMs. Constrained decoding for AR LLMs compiles grammars into incremental automata that mask invalid logits at each step [Willard and Louf, 2023, Dong et al., 2024, Microsoft, 2025], with much recent work spent on making that mask cheap [Ugare et al., 2025, Park et al., 2025] or on extending it to semantic [Poesia et al., 2022] and type-system [Mündler et al., 2025b] constraints. Geng et al. [2023] give the framing we adopt, that a formal grammar describes the output space of structured NLP tasks in general rather than of serialisation formats alone, and Tuccio et al. [2025] carry it into open-ended generation; our constraint is context-free for the same reason, although we evaluate on JSON Schema. All of it assumes monotonic left-to-right generation, so out-of-order commitment never arises, and that is the central problem in dLLMs.

Constrained decoding for dLLMs. Mündler et al. [2025a] propose a CFG-constrained dLLM decoder, intersecting the grammar with an automaton over the generated prefix. DINGO [Suresh et al., 2025] performs MAP inference over a deterministic finite automaton weighted by the model’s mean-field predictions, and Dang and Ermon [2026] generalise this to exact sampling from the constrained mean-field posterior under arbitrary remasking schedules, restoring a hard guarantee. Both require the constraint as a finite automaton, which bounds them to regular languages; we trade that guarantee for a context-free constraint language and a schedule-agnostic layer. Cardei et al. [2025] also prevent rather than repair, by a differentiable projection inside the sampling loop, which reaches constraints admitting a continuous relaxation but not grammar membership, a discrete predicate a parser handles directly. LAVE [Zhang et al., 2026] verifies each frontier token by sampling K random completions, accepting only if one is grammar-valid; it is the strongest published baseline we can run, and Nakshatri et al. [2025] make the same K roll-outs cheaper in the AR setting with a draft model. Plan-style decoders [Li et al., 2026] emphasise structure-aware planning.

¹On JSONSchemaBench medium an automata toolchain compiles only 383 of 586 schemas where an incremental context-free parser handles 511 (Appendix H).

²Against a vocabulary of 126,349, a live parser state admits a median of just 106 tokens (§3.4, Appendix A).

85 **Viterbi decoding for non-AR models.** Lattice search is well established in non-autoregressive
 86 generation [Deng and Rush, 2021]: Shao et al. [2022] apply Viterbi decoding over Directed Acyclic
 87 Transformers and Control-DAG [Chen et al., 2024] enforces lexical constraints with weighted finite-
 88 state automata. Both search a model-defined lattice in which output length must be constrained, or an
 89 unnormalised likelihood favours trivially short paths. Ours is defined by the parser and bounded by
 90 the violated span, so length is fixed rather than searched and the length-control term those methods
 91 need has nothing to correct, which §4.4 measures.

92 3 Method

93 3.1 Problem statement

94 Let G be a grammar and let its *incremental parser* be an object that processes tokens left to right.
 95 After accepting a prefix, the parser maintains a *state* q defining the valid continuations permitted
 96 by G . Write $\mathcal{L}(q) \subseteq V$ for the set of tokens the parser can consume from q , its outgoing edge set
 97 (Appendix F shows the concrete interface). The decoder holds a sequence x over $V \cup \{\text{MASK}\}$; its
 98 frontier c is the length of the longest prefix the parser has consumed, and q_c the state it reached there.
 99 Because unmasking follows confidence rather than position, the region beyond the frontier mixes
 100 masked slots with tokens the model has already placed:

$$\underbrace{x_0 \cdots x_{c-1}}_{\text{consumed, parser at } q_c} \quad \bigg| \quad \underbrace{x_c}_{\text{frontier}} \quad \bigg| \quad \underbrace{\text{MASK } x_{c+2} \text{ MASK } x_{c+4} \cdots}_{\text{placed but never parsed}}$$

101 Only x_c is under the parser’s control: frontier masking guarantees $x_c \in \mathcal{L}(q_c)$ when x_c is chosen this
 102 step, while every position to its right was sampled with no constraint at all.

103 Those positions are not provisional. Masked diffusion is *carry-over unmasking* [Sahoo et al., 2024,
 104 Shi et al., 2024]: the reverse posterior copies any already-revealed position unchanged, so a step only
 105 ever replaces MASK symbols, which is what LLaDA’s reference sampler implements. (Samplers that
 106 deliberately remask revealed positions relax the property; we do not assume one.) Tokens placed
 107 beyond the frontier are therefore immutable commitments. As the remaining masks fill in, the parser
 108 advances across the newly contiguous region and may reach a position $v > c$ where $x_v \notin \mathcal{L}(q_v)$: a
 109 *violation*, a final token the grammar refuses.

110 **Why this is challenging.** Every constrained decoder for an autoregressive model is a filter: it
 111 inspects one position’s distribution and removes what the grammar forbids. A violator sits past the
 112 frontier, where nothing can be filtered, only overwritten. Local repair rarely suffices, because the
 113 tokens that interact to cause the violation are often several positions away. The problem is a *search*
 114 over joint assignments to a region, not a filter over one distribution.

115 3.2 Repair as bounded search

116 Given a violator v , the layer repairs a bounded span $[v, e)$: it rewrites committed positions and never
 117 fills holes. Three things bound e . A forward scan from v reads the tokens already in place and stops
 118 at the first one the grammar accepts *at bracket depth zero*, which is the span end when it finds one;
 119 failing that it stops at a compute budget of $L = 48$ positions. The span is then truncated at the first
 120 MASK, since a hole is the sampler’s to fill and not ours. The depth test is the part that matters: without
 121 it the span can end inside an unclosed array or object, where the cheapest legal continuation is the
 122 closing bracket, and the joint solve then collapses the rest of the structure to $[]$ or $\{\}$, the degeneracy
 123 §4.2 measures with the content floor. Bracket depth is the JSON instance of a general requirement,
 124 that a span end where the grammar can resynchronise; another constraint family would need its own
 125 test. Appendix A reports which of the three binds in practice, and how long the resulting spans are.

126 Over that span we seek the assignment y maximising $\sum_p \log p_\theta(y_p \mid x)$ subject to the parser
 127 accepting $y_v \cdots y_{e-1}$ from q_v . §3.3 solves it as a Viterbi pass and §3.4 says where the candidates at
 128 each position come from.

129 A rejection at v enters the cascade of Figure 1. Termination is tested first: EOS is not consumable,
 130 so a finished document reaches the frontier looking exactly like a violation. We separate the
 131 two on the grammar rather than on a step budget. The harness ends the document only when




Figure 1: Where the joint repair fires. The row is tried left to right and each stage is reached only when the one before it did not settle the position; every outcome but emission returns to the next denoising step. The shaded box is this paper.

132 $\mathcal{L}(q) \setminus \{\text{EOS}, \text{EOT}\} = \emptyset$ at an accepting q , where ending is not a choice anyone is making; the model
 133 ends it when it places a stop token at the frontier and q is accepting, where the model has asked and
 134 the grammar need only permit. Next is a bounded single-token retry, which walks the model’s ranked
 135 logits at v for up to 10 candidates and takes the first the parser accepts. That is filter-then-propose,
 136 the pattern §3.4 argues against, but a miss here is free because the position falls through to the joint
 137 search, and it resolves 72% of violations at negligible cost. Only then is the span solved jointly (§3.3);
 138 if no assignment exists, v is remasked for the sampler to re-derive.

139 3.3 Joint repair by Viterbi over parser states

140 Algorithm 1 gives the search. It is a Viterbi pass over layers indexed by position, where a node is a
 141 reachable *parser state* rather than a token history. Candidates that leave the parser in the same state
 142 κ' collapse to the single highest-scoring path (line 11), and that merge is what bounds the search.
 143 Enumerating assignments over the span keeps $O(k^{\text{span}})$ hypotheses alive; merging caps each layer at
 144 the number of distinct keys $|S|$, giving $O(\text{span} \cdot |S| \cdot k)$ parser operations and $O(|S|)$ live parsers.
 145 The key $\kappa(q) = \text{bytes}(\mathcal{L}(q))$ is the parser’s own mask, which line 5 already computed to build C ,
 146 so the merge costs no parser call of its own. Nothing bounds $|S|$ a priori, it being the number of
 147 distinct masks the grammar can present at one position, but it stays in the single digits across our
 148 runs, and span length is no longer an exponent whatever it is.

149 **What merging costs.** κ is a state *abstraction*, not a state identifier: two states can admit the
 150 same tokens and still differ in what the grammar demands later, so merging on it is not exact. It is
 151 nonetheless sound, because the key prunes but never reconstructs. Each retained node holds the parser
 152 it reached and the backtrace follows predecessor links, so what the pass returns is always a sequence
 153 some parser consumed from q_v , and a collision cannot produce a span the grammar rejects. What a
 154 collision can cost is the objective: the lower-scoring of two genuinely different states is dropped, and
 155 where the survivor dead-ends first the span ends early, which Algorithm 1 treats as finishing rather
 156 than failing. Appendix F measures that cost by widening the merge, and finds it changes the returned
 157 assignment at 0.33% of solves. Appendix E states separately what the search guarantees, what it
 158 assumes, and what each approximation costs.

159 3.4 Candidates from the automaton

160 At each lattice node the search needs a set of tokens to try. We draw it from the parser rather than
 161 from the model: for a live state q we read $\mathcal{L}(q)$, the tokens that state can consume, and rank *within* it
 162 by model likelihood,

$$\mathcal{C}(q, p) = \text{top-}k_{t \in \mathcal{L}(q) \setminus \{\text{EOS}\}} \log p_{\theta}(t \mid x, p). \quad (1)$$

163 The alternative is to take the model’s top- k predictions over V and retain only those the parser accepts.
 164 Filtering after sampling has a structural failure mode: $|\mathcal{L}(q)|$ is set by the grammar rather than by a
 165 hyperparameter, so no choice of k removes the risk: where the admitted set is small the draw omits it
 166 entirely and the DP reports “no path” at a position where one exists (Appendix A works one through).

167 Drawing from $\mathcal{L}(q)$ removes that failure mode by construction, because the proposal *is* the accepted
 168 set. What k bounds is therefore likelihood, not validity. Validity is unconditional: every candidate
 169 comes from $\mathcal{L}(q)$, so no k , however small, lets the search emit an invalid span. An answer is always
 170 returned, since no live state we measured has an empty legal set, and C is empty only where the

Algorithm 1 Joint repair of a violated span

Require: violator v , span end e , parser state q_v , log-probabilities $\log p_\theta$, candidate budget k
Ensure: an assignment over $[v, e']$, $e' \leq e$, that the parser accepts from q_v

```
1:  $S \leftarrow \{ \kappa(q_v) \mapsto (q_v, 0) \}; \quad B \leftarrow \emptyset$  ▷ key  $\mapsto$  (parser, score); back-pointers
2: for  $p = v$  to  $e - 1$  do
3:    $W \leftarrow \emptyset$  ▷ best incoming path per successor state
4:   for all  $(\kappa, (q, s)) \in S$  do
5:      $C \leftarrow$  the  $k$  most probable tokens of  $\mathcal{L}(q) \setminus \{\text{EOS}\}$  ▷ §3.4
6:      $C \leftarrow C \cup \{x_p\}$  if  $x_p \in \mathcal{L}(q)$  ▷ keep the model's own token reachable
7:     for all  $t \in C$  do
8:       advance  $q$  by  $t$ ;  $\kappa' \leftarrow \kappa(q)$ ;  $s' \leftarrow s + \log p_\theta(t)_p$ 
9:       rollback  $q$  one token ▷ probe without copying
10:      if  $\kappa' \notin W$  or  $s' > \text{score}(W[\kappa'])$  then
11:         $W[\kappa'] \leftarrow (\kappa, q, t, s')$  ▷ merge: paths reaching  $\kappa'$  collapse to the best one
12:      end if
13:    end for
14:  end for
15:  if  $W = \emptyset$  then
16:    break ▷ nothing consumable: the span ends here, it does not fail
17:  end if
18:   $S \leftarrow \{ \kappa' \mapsto (\text{clone}(q) \text{ advanced by } t, s') \}$  and  $B[p, \kappa'] \leftarrow (\kappa, t)$ , for each  $(\kappa', (\kappa, q, t, s')) \in W$ 
19: end for
20: return the tokens on the path  $B$  traces back from  $\arg \max_{\kappa \in S} \text{score}(\kappa)$ 
```

171 grammar admits nothing but a stop token, in which case the document is finished and there is nothing
172 to place. Only optimality is relative to the lattice: the pass returns the best assignment the candidate
173 sets span, which is the best the grammar accepts exactly where those sets are complete, and where
174 $|\mathcal{L}(q)| \leq k$ the edge set is enumerated rather than sampled. Appendix A gives the distribution of
175 $|\mathcal{L}(q)|$, how often it fits inside k , and where the span budget L binds.

176 4 Experiments

177 4.1 Setup and metrics

178 **Datasets.** JSONSchemaBench [Geng et al., 2025] supplies JSON schemas without reference
179 answers, in three difficulty splits of 577, 586 and 368 instances. LLGUIDANCE implements a large
180 subset of JSON Schema, and a schema it cannot compile is excluded from every arm since no method
181 runs on it: 19, 75 and 99 respectively, leaving $n = 558, 511, 269$. JSON-Mode-Eval supplies prompts
182 paired with a schema *and* a reference completion; we use the 272-instance extension [ETH SRI,
183 2025] of the original 100-row release [Nous Research, 2024], as released with CD-CFG.

184 **Models.** All arms run LLaDA-8B-Instruct [Nie et al., 2025] with $T=128$ denoising steps and block
185 length 32, on a single A100 per shard, with grammars compiled by llguidance [Microsoft, 2025]
186 except where noted. Appendix D lists the remaining constants.

187 **Baselines.** *Vanilla* is the same decoder with no constraint of any kind, and marks what the backbone
188 produces unaided. *CD-CFG* [Mündler et al., 2025a] is run from the authors' released implementation,
189 which builds its constraint with RUSTFORMLANG, the formal-language library that ships with it, and
190 scored with that code's `autocomplete_valid` post-processing, the configuration its authors ship;
191 §H treats the coverage difference separately and reports what the post-processing is worth. *LAVE*
192 [Zhang et al., 2026] is likewise run from its authors' released code, instrumented only to record
193 per-operation timing, under a 120 second per-instance cap that applies to every arm and scores an
194 instance reaching it as a failure. *no-DP* is our own system with the repair layer switched off: *the*
195 *same binary, the same checkpoint, the same schedule and the same seed*. Upon a violation, it retains
196 a bounded single-token retry and, if that is exhausted, defers the position back to the sampler. Both
197 arms share the system-level optimisations that keep the decoder fast, which Appendix C describes.

Table 1: The difficulty gradient on LLaDA-8B-Instruct. Each arm is scored on the instances its own toolchain compiles, so n differs: LAVE, no-DP and DPG share the LLGUIDANCE set exactly, while $\dagger$ marks CD-CFG, whose RUSTFORMLANG front-end accepts a smaller subset (Appendix H). Times are in seconds and TO counts the 120 s cap. rs/inst counts positions handed back to the sampler per instance; its unit differs by arm (remasks, lookahead draws, rejected candidates), so only no-DP and DPG compare there. It does not count tokens the joint search rewrites (§4.2). DPG is DPGrammar.

Split	Method	n	schema@1 $\uparrow$	content $\uparrow$	mean $\downarrow$	p50 $\downarrow$	p95 $\downarrow$	p99 $\downarrow$	TO $\downarrow$	rs/inst $\downarrow$
easy	Vanilla	577	62.9	30.2	8.34	8.37	11.23	14.98	0	—
	CD-CFG $\dagger$	449	79.1	36.3	6.22	6.72	11.54	15.98	0	49.2
	LAVE	558	80.1	40.5	30.63	6.71	120.01	120.01	94	207.3
	no-DP (ours)	558	97.5	45.2	7.58	7.58	11.47	15.09	0	20.7
	DPG (ours)	558	97.7	49.3	5.31	3.40	15.11	29.06	0	0.6
medium	Vanilla	586	49.8	47.1	14.89	14.36	21.82	31.67	0	—
	CD-CFG $\dagger$	383	64.8	58.0	22.71	14.16	112.24	120.00	19	57.3
	LAVE	511	78.3	74.0	38.50	25.70	120.00	120.01	60	253.5
	no-DP (ours)	511	91.8	74.6	13.45	13.40	21.43	29.99	0	33.2
	DPG (ours)	511	96.7	84.9	16.75	15.05	36.21	63.24	0	1.2
hard	Vanilla	368	20.7	20.4	44.64	35.43	103.48	120.00	4	—
	CD-CFG $\dagger$	245	32.2	29.8	57.61	41.53	120.00	120.15	72	50.2
	LAVE	269	39.8	39.4	40.47	38.35	120.00	120.00	30	115.8
	no-DP (ours)	269	74.3	53.9	31.41	26.69	61.91	101.51	0	56.0
	DPG (ours)	269	87.7	71.7	42.66	35.24	96.58	115.96	2	1.5

198 The closest work we do *not* run is DINGO [Suresh et al., 2025] and the exact sampling of Dang and
199 Ermon [2026]: both need the constraint as a finite automaton, which the recursive schemas here do
200 not admit, and neither had a released implementation at the time of writing. §4.3 reports what that
201 guarantee would have traded, measured with our own implementation of exact constrained inference.

202 **Metrics.** Three quality metrics. *schema@1* is validation against the instance’s own schema under
203 *jsonschema*. *content* additionally requires five *populated leaves*, the scalar values anywhere in the
204 document, since validity alone cannot separate an empty structure from a populated one; Table 2
205 sweeps the threshold and Appendix K shows a pair the two order differently. *functional@1* requires a
206 match against the reference completion, so it is JSON-Mode-Eval only. Two cost metrics: hand-backs
207 per instance (*rs/inst*, defined with Table 1) and wall-clock latency, reported as mean, p50, p95 and
208 p99 with timeouts at a 120 s cap.

209 4.2 Main results

Table 2: JSONSchemaBench medium ($n=511$) as the content floor rises. Each row requires validity *and* at least κ populated leaves.

	LAVE	no-DP	DPGrammar	Δ vs no-DP	Δ vs LAVE
validity only	78.3	91.8	96.7	+4.9	+18.4
$\wedge \kappa=1$	78.3	84.9	93.7	+8.8	+15.4
$\wedge \kappa=3$	76.7	79.1	90.6	+11.5	+13.9
$\wedge \kappa=5$	74.0	74.6	84.9	+10.3	+10.9
$\wedge \kappa=10$	56.0	50.7	60.5	+9.8	+4.5
$\wedge \kappa=15$	33.7	29.5	35.0	+5.5	+1.3

210 Repair pays in proportion to how jointly the grammar constrains a region, and schema validity hides
211 most of what it pays. On medium the layer lifts validity by 4.9 pp over the unrepaired baseline while
212 cutting hand-backs by 28 $\times$, from 33.2 to 1.2 per instance. Hand-backs are not the whole of the
213 corrective work: in their place the layer runs 0.62 joint solves per instance and rewrites 3.0 tokens,
214 so the positions it touches fall from 33.2 to 4.2: eight rather than twenty-eight. Two trends say
215 more than the headline. Raising the content floor doubles the margin, from +4.9 to +10.3 pp at
216 five populated leaves and +11.5 at three (Table 2), because single-token repair takes the cheapest

217 legal token at a violation and a closing bracket usually is one: the document collapses exactly where
 218 the joint solve would have continued it. Raising the difficulty does the same, by +0.2, +4.9 and
 219 +13.4 pp of validity and +4.1, +10.3 and +17.8 pp of content across *easy*, *medium* and *hard*, the
 220 first a single instance, which is what joint repair predicts: violations become more joint as schemas
 221 acquire nesting, required-property ordering and cross-field constraints. Against LAVE the shape
 222 inverts, narrowing from +18.4 to +1.3 pp as the floor rises, because LAVE never rewrites and its
 223 surviving documents were never at risk of collapsing. Against LAVE we buy frequency; against the
 224 unrepaired baseline we buy content (Appendix K).

225 That margin is not bought with time, and it does not depend on which instances each arm is scored on.
 226 The constraint machinery costs what no-DP’s does, so the 3.3 s is not the search but the decoding the
 227 repair makes possible: a repaired document keeps going where no-DP stops, and where it fires the
 228 total rises by 96% but the cost per token emitted by only 15% (Appendix J). That brevity is premature
 229 termination rather than economy: 18.8% of no-DP’s valid *medium* outputs carry fewer than five
 230 leaves against 12.1% of ours. On *easy*, where repair is rarely needed, DPGrammar is the faster arm.
 231 The tail does grow, from 21.4 to 36.2 s at p95, and two *hard* instances reach the cap where no-DP
 232 reaches none. The ordering is unchanged on a shared population, whether the 351 schemas both
 233 toolchains compile or the 451 instances LAVE finishes within the cap (Appendix H).

234 The content margin survives a change of schedule and grows where the schedule is harder: +13.1 pp
 235 when the whole sequence is committed in one block rather than in blocks of 32, and +17.2 when
 236 positions are revealed in random order rather than by confidence, against +10.3 here. The validity
 237 margin travels less evenly, from +15.1 under the first to +0.4 under the second, and it tracks the step
 238 budget, not the search. Random order produces 2.3× the violations, repairing them costs denoising
 239 steps, and 46% of instances exhaust the 128-step schedule before the document closes, 98% of what
 240 DPGrammar fails on there; the search itself dead-ends once in 733 calls (Appendix I).

241 Not every residual failure is a failure of the search. LLGUIDANCE pins an object’s properties to their
 242 declaration order, so a model that writes them in another order produces a violation *jsonschema*
 243 would not call an error at all, and whose correct repair is a movement rather than a substitution,
 244 which a positional DP cannot express. This accounts for 6.6% of the 1,106 violations the reported
 245 run records, and for 50% of the 120 that arise where the grammar admits a single token, the states in
 246 which the search has no alternative to write (Appendix B). It is a limit of overwriting in place rather
 247 than of solving the span jointly.

248 4.3 Two comparisons the main table cannot make

Table 3: JSON-Mode-Eval, $n=272$, LLaDA-8B-Instruct. *schema@1* is the same quantity as in Table 1; *functional@1* additionally requires the output to match the reference completion.

	Vanilla	CD-CFG	LAVE	no-DP	DPGrammar
<i>schema@1</i> ↑	85.3	97.8	98.2	97.4	98.2
<i>functional@1</i> ↑	48.5	52.9	53.7	53.7	52.6

249 Constraining structure does not buy semantics. JSON-Mode-Eval is where the two separate, since it
 250 pairs each prompt with a reference completion: every constrained arm gains 12 to 13 pp of schema
 251 validity over Vanilla while *functional@1* stays inside 5.2 pp. The layer has little to do here: no-DP
 252 already validates 97.4% of these instances, so the joint search fires on 14 of 272 and the two arms
 253 agree byte for byte on 256. What separates their *functional@1* is five instances, one in our favour
 254 and four against, which an exact McNemar test places inside noise ($p = 0.38$). On those 14 the layer
 255 takes *schema@1* from 12 to 14, rewriting 5.1 tokens over 16 solves, and moves *functional@1* by one
 256 win against two losses. The gap measures this sample, not what rewriting costs an answer; §4.2 runs
 257 the layer under joint-constraint load.

258 **What prevention would have cost.** Prevention buys its guarantee three times over. Run on the
 259 same backbone and schedule, exact inference over the constrained posterior reaches a quarter fewer
 260 schemas than the incremental parser, is three to five times slower, and returns a median of one scalar
 261 leaf against our thirteen. That last is the content floor from the other side: the smallest document a
 262 schema admits is both valid and cheap, so a decoder optimising probability under a hard constraint is
 263 drawn to it (Appendix G).

264 4.4 Ablations: what moves the metric is the search space, not the objective

Table 4: Ablation of the repair layer’s components, on JSONSchemaBench medium with one seed. The upper block changes what the search ranges over; the merge width is measured on the search itself rather than end to end (§3.3). The lower block changes the objective or the admission rule and leaves schema@1 where it is. One-char string leaves are the share of an output’s string leaves that hold a single character, averaged over outputs.

Component	Varied	Measured effect
Candidate source	vocabulary top- $k \rightarrow$ automaton	schema@1 95.5 $\rightarrow$ 96.7; dead ends 23.3 $\rightarrow$ 1.0%
Repair window	single token $\rightarrow$ full span	one-char string leaves 12.3 $\rightarrow$ 7.9%
Paths kept per key	1 $\rightarrow$ 4	assignment changes at 0.33% of solves
Objective	$\log p \rightarrow$ min-edit	−0.5 edits, 93.9% identical
Candidate admission	always admit committed token	none (identical output)
Stop tokens	excluded $\rightarrow$ kept in $\mathcal{C}$	99.6% identical; 48× futile probes
Post-hoc enrichment	on $\rightarrow$ off	none (identical output)
Edit penalty	$\lambda \in \{0, 2\}$	none (never non-zero)

265 We vary one component at a time, holding backbone, seed and schedule fixed, so what follows is
 266 evidence about this configuration, not a law. Three effects the table cannot show are worth stating.
 267 Filtering the vocabulary after sampling abandons 273 positions the edge set recovers, and the three
 268 dead ends it does produce are not proposal misses but states admitting nothing but a stop token,
 269 the one case $\mathcal{C}$ is empty by design. Minimum-edit attributes 90% of what the DP rewrites to the
 270 grammar rather than to the objective. And the edit penalty stays at zero although Viterbi decoders
 271 for non-autoregressive models need it to stop an unnormalised objective collapsing toward short
 272 paths [Shao et al., 2022, Chen et al., 2024]: our span is fixed by the violation rather than chosen by
 273 the search, so there is no collapse to prevent.

274 5 Conclusion

275 We present DPGrammar, a repair layer for grammar-constrained decoding in diffusion language
 276 models, which commit tokens out of order and can therefore only repair a violation after the fact.
 277 Casting that repair as a dynamic program over parser states lets a violated region be decided as a
 278 whole rather than one position at a time, with paths that leave the parser in the same state merging
 279 into one. Because candidates are read off the parser rather than the vocabulary, bounding that search
 280 trades likelihood and never validity. What that buys is not a constant offset. The margin over repairing
 281 one position at a time doubles once outputs must carry content, and grows monotonically with how
 282 jointly a schema constrains the region it fails in, which is where prevention pays instead: in coverage,
 283 in latency, and in documents that satisfy a schema by saying nothing. Validity is held at the smallest
 284 intervention the grammar allows, by a search that fires only after a rejection, proposes only tokens
 285 the parser admits, follows the model’s ranking wherever the grammar leaves a choice, and overwrites
 286 a position only where it does not.

287 6 Limitations and future work

288 Our results are one backbone, one seed and one constraint family, and the last of those bounds more
 289 than the numbers: a constraint whose states are harder to tell apart by their legal-token sets would pay
 290 more for the state key than the 0.33% we measure here (§3.3). The layer also earns nothing where
 291 violations are local rather than joint: across 64 SMILES instances the single-token retry absorbed all
 292 9 violations the parser raised, so the joint search never ran; that is the difficulty gradient read at its
 293 other end, but it leaves the boundary between those grammars and JSON Schema unmapped, and
 294 other backbones and constraint families would place it.

295 A more fundamental direction is the parser. Ours is written for left-to-right decoding, which is why
 296 the constraint can be enforced at one frontier and everything past it only repaired. A checker that
 297 tracked several disconnected regions at once would let the grammar be enforced wherever a step
 298 touches, and give a joint repair more of the sequence to work with.

References

- Michael Cardei, Jacob K. Christopher, Thomas Hartvigsen, Bhavya Kailkhura, and Ferdinando Fioretto. Constrained discrete diffusion. In *Advances in Neural Information Processing Systems*, 2025. arXiv:2503.09790.
- Jinghong Chen, Weizhe Lin, Jingbiao Mei, and Bill Byrne. Control-DAG: Constrained decoding for non-autoregressive directed acyclic T5 using weighted finite state automata. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Short Papers)*, 2024.
- Meihua Dang and Stefano Ermon. Constrained decoding for diffusion language models via efficient inference over finite automata. *arXiv preprint arXiv:2607.07026*, 2026.
- Yuntian Deng and Alexander M. Rush. Sequence-to-lattice models for fast translation. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, Punta Cana, Dominican Republic, 2021. URL <https://aclanthology.org/2021.findings-emnlp.318/>.
- Yixin Dong, Charlie F Ruan, Yaxing Cai, Ruihang Lai, Ziyi Xu, Yilong Zhao, and Tianqi Chen. Xgrammar: Flexible and efficient structured generation engine for large language models. *Proceedings of Machine Learning and Systems* 7, 2024.
- ETH SRI. json-mode-eval-extended. Hugging Face dataset, 2025. URL <https://huggingface.co/datasets/eth-sri/json-mode-eval-extended>. Extension of NousResearch/json-mode-eval to 272 instances.
- Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. Grammar-constrained decoding for structured NLP tasks without finetuning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 10932–10952, Singapore, 2023. URL <https://aclanthology.org/2023.emnlp-main.674/>.
- Saibo Geng, Hudson Cooper, Michał Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. Jonschemabench: A rigorous benchmark of structured outputs for language models, 2025. URL <https://arxiv.org/abs/2501.10868>.
- Guidance AI. Guidance: A guidance language for controlling large language models, 2023. URL <https://github.com/guidance-ai/guidance>.
- Miao Li, Hanyang Jiang, Sikai Cheng, Hengyu Fu, Yuhang Cai, Baihe Huang, Tinghan Ye, Xuanzhou Chen, and Pascal Van Hentenryck. Plan, verify and fill: A structured parallel decoding approach for diffusion language models, 2026. URL <https://arxiv.org/abs/2601.12247>.
- Microsoft. LLMGuidance: Fast structured outputs for LLMs. GitHub repository, 2025. URL <https://github.com/guidance-ai/llguidance>.
- Niels Mündler, Jingxuan He, Hao Wang, Koushik Sen, Dawn Song, and Martin Vechev. Type-constrained code generation with language models. *Proceedings of the ACM on Programming Languages*, 9(PLDI), 2025b. doi: 10.1145/3729274.
- Niels Mündler, Jasper Dekoninck, and Martin Vechev. Constrained decoding of diffusion llms with context-free grammars, 2025a. URL <https://arxiv.org/abs/2508.10111>. arXiv:2508.10111.
- Nishanth Sridhar Nakshatri, Shamik Roy, Rajarshi Das, Suthee Chaidaroon, Leonid Boytsov, and Rashmi Gangadharaiyah. Constrained decoding with speculative lookaheads. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 4681–4700, 2025. URL <https://aclanthology.org/2025.naacl-long.239/>.
- Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models, 2025. URL <https://arxiv.org/abs/2502.09992>.

346 Nous Research. json-mode-eval. Hugging Face Datasets, 2024. URL <https://huggingface.co/datasets/NousResearch/json-mode-eval>.
347

348 Kanghee Park, Timothy Zhou, and Loris D’Antoni. Flexible and efficient grammar-constrained
349 decoding. In *International Conference on Machine Learning*, 2025. arXiv:2502.05111.

350 Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and
351 Joseph E. Gonzalez. The Berkeley function calling leaderboard (BFCL): From tool use to agentic
352 evaluation of large language models. In *International Conference on Machine Learning*, 2025.

353 Gabriel Poesia, Alex Polozov, Vu Le, Ashish Tiwari, Gustavo Soares, Christopher Meek, and
354 Sumit Gulwani. Synchronesh: Reliable code generation from pre-trained language models. In
355 *International Conference on Learning Representations*, 2022.

356 Subham Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin Chiu,
357 Alexander Rush, and Volodymyr Kuleshov. Simple and effective masked diffusion language
358 models. In *Advances in Neural Information Processing Systems*, volume 37, pages 130136–
359 130184, 2024.

360 Chenze Shao, Zhengrui Ma, and Yang Feng. Viterbi decoding of directed acyclic transformer for non-
361 autoregressive machine translation. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 4390–4397, 2022.
362

363 Jiaxin Shi, Kehang Han, Zhe Wang, Arnaud Doucet, and Michalis Titsias. Simplified and generalized
364 masked diffusion for discrete data. In *Advances in Neural Information Processing Systems*,
365 volume 37, pages 103131–103167, 2024.

366 Tarun Suresh, Debangshu Banerjee, Shubham Ugare, Sasa Misailovic, and Gagandeep Singh. Dingo:
367 Constrained inference for diffusion llms, 2025. URL <https://arxiv.org/abs/2505.23061>.

368 Gabriele Tuccio, Luana Bulla, Maria Madonia, Aldo Gangemi, and Misael Mongiovi. GRAMMAR-
369 LLM: Grammar-constrained natural language generation. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 3412–3422, 2025. URL <https://aclanthology.org/2025.findings-acl.177/>.
371

372 Shubham Ugare, Tarun Suresh, Hangoo Kang, Sasa Misailovic, and Gagandeep Singh. SynCode:
373 LLM generation with grammar augmentation. *Transactions on Machine Learning Research*, 2025.
374 arXiv:2403.01632.

375 Brandon T Willard and Rémi Louf. Efficient guided generation for large language models. *arXiv preprint arXiv:2307.09702*, 2023.
376

377 Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui Wu, Xin Jiang, Zhenguo Li, and Lingpeng
378 Kong. Dream 7b: Diffusion large language models, 2025. URL <https://arxiv.org/abs/2508.15487>.
379

380 Yitong Zhang, Yongmin Li, Yuetong Liu, Jia Li, Xiaoran Jia, Zherui Li, and Ge Li. Lookahead-
381 then-verify: Reliable constrained decoding for diffusion LLMs under context-free grammars.
382 *Proceedings of the ACM on Software Engineering*, 3(ISSTA), 2026.

383 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao,
384 Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang:
385 Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, 2024.
386

A The distribution of legal token sets

Section 3.4 rests on a claim about how many tokens an incremental parser actually admits at a live state. This appendix gives the measurement and the full distribution.

Procedure. We take the first 40 outputs, in file order, of the DPGrammar run on JSONSchemaBench medium (seed 0, $T=128$) whose schemas compile, and replay each one through `llguidance` 1.7.0 under teacher forcing: the parser is initialised from the instance’s schema, and the tokens the model actually produced are fed to it one at a time. Before consuming each token we call `compute_logit_bias()` and count the non-zero entries of the returned mask; that count is $|\mathcal{L}(q)|$, the number of tokens the parser would accept next. Replay stops at the first token the parser rejects, so every state we record is one the decoder genuinely reached. This yields 6,657 live states over a vocabulary of $|V| = 126,349$. The measurement is offline and model-free, which needs only the schema and the emitted text, so it can be reproduced without a GPU.

The distribution. Figure 2 shows the result. The distribution is not merely skewed, it is bimodal: 38.4% of states admit fewer than 100 tokens and 29.7% admit more than 10,000, leaving the decade between 10^4 and 10^5 nearly empty. The median is 106, and the extremes are sharper still: before its first token a JSON object grammar admits 2 of 126,349.




Figure 2: How many tokens an incremental parser admits, over 6,657 live states on JSONSchemaBench medium ($|V| = 126,349$).

What the order of proposal and filter costs. The distribution above is why §3.4 draws candidates from $\mathcal{L}(q)$ rather than ranking V and filtering: at the 38% of live states admitting fewer than 100 tokens, a top-100 draw is aiming at a small target. Reversing the order in our own implementation costs 1.2 pp of schema@1, and dead-ends the search at 93 of its 399 repair calls against 3 of 315 (§4.4). Those 93 are the failure this ordering cannot design away: the draw held no token the parser would accept, so the search reported no path where one existed and handed the position back, 273 positions in all.

One site, both orders. Instance o65372 gives the failure in miniature. At position 17 the model has written " , and the parser refuses it: the schema is spelling out a mandated property name, and from that state it admits *one* token of the 126,349 in the vocabulary, the next character of that name. The model ranks the vocabulary without seeing the grammar, so its top-100 has no particular reason to contain that character, and here it does not. Filtering after sampling therefore leaves nothing: the candidate set is empty at the first layer, the DP dead-ends before it places anything, and position 17 goes back to the sampler. Reading the edge set first leaves the same single token, which is all the search needs; from it the DP repairs a seven-position span in three edits. Nothing about the grammar, the model or the state differs between the two runs. Only the order in which they are consulted does.

Where the two modes come from. The obvious explanation, that the tight mode is structural punctuation and the wide mode is string content, is only half right. Structural positions (3,429 states)

are uniformly tight, with a median of 104 and a quartile range of 95 to 324. The bimodality lives entirely inside strings (3,228 states, $p_{25} = 3$, median = 125,308), where it separates two kinds of string: a free-form value admits almost the whole vocabulary, whereas a partially spelled property name admits a handful of continuations. Both modes matter for repair, and violations do not fall evenly across them. Of the 1,106 the reported run records, 43.9% occur at a state admitting fewer than 100 tokens against 38.4% of live states overall, and 22.4% at one admitting more than 10,000 against 29.7%: the tight mode is over-represented, though only modestly. Its tail is not modest. One violation in nine, 120 of 1,106, arises where the grammar admits a single token, and those are the states a top- k draw over the vocabulary has no reason to hit.

When the automaton gives the complete candidate set. Where $|\mathcal{L}(q)| \leq k$, Eq. 1 enumerates the edge set exactly rather than sampling it, and the candidate set is provably complete:

k	8	32	64	100	128	256
states with $ \mathcal{L}(q) \leq k$	16.4%	17.9%	18.1%	39.7%	54.0%	54.6%

At the reported $k=100$ this holds at 39.7% of live states. The tight mode clusters just above 100, so $k=128$ reaches 54.0% and $k=256$ only 0.6 beyond that: the mode is exhausted by 128. Whether that extra coverage is worth anything is a question we can answer with the same replay used for the merge width. Solving each site twice, at $k=100$ and at $k=128$ and otherwise identically, the two returned the same assignment at every one of 320 solves over 161 states, and not because the budget was idle: 60% of those states admit more than 100 tokens and 31% more than 128. What makes k a safe knob is in any case not its value but the order it is applied in, since a smaller k costs likelihood and cannot cost validity (§3.4). The fraction varies widely across schemas (at $k=32$: median 18.4%, interquartile range 9.8–28.3%, and 100% for the most constrained schema in the sample), so it is a property of the schema rather than a constant of the tokenizer.

The span budget. L bounds the forward scan of §3.2 rather than the candidate set. Over the 315 repair sites of the reported run the scan ends at the first MASK at 75.6% of them and at a depth-zero junction at 9.8%, which leaves the budget itself binding at 7.9%; the remaining 6.7% end where the probe can no longer advance. Spans run to 13.8 positions on average and 9 at the median. A smaller budget would bind more often, at 13.3% of sites for $L=32$, 19.0% for 24, 31.7% for 16 and 52.4% for 8, so $L=48$ sits where that curve has flattened and a span it truncates is already an outlier.

Two facts the method depends on. First, no live state had an empty legal set (0 of 6,657). This is what makes the DP total: a live state always has a continuation, so the candidate set is never empty and the search can always complete the span. Our termination rule (§3.2) treats an empty mask as a failed read rather than as a stop signal for this reason. Second, the mask read is cheap (0.025 ms mean, 0.012 ms median) and the implementation already performs it to compute the state key used for path merging, so drawing candidates from the automaton adds no parser calls beyond those already paid for.

B Ordering artifacts

Ilguidance compiles a JSON schema into a grammar that fixes the order of an object’s properties to their declaration order: $\{ "a" : 1, "b" : 2 \}$ is rejected when the schema declares b first, although both orders satisfy the schema and `jsonschema` accepts either. A model that writes the properties in a different order therefore produces a violation that is not an error, and the correct repair, moving a block, is not expressible by a positional DP, which can only overwrite in place. The repair then writes the demanded key name over a value intended for another field. We measured this class at 6.6% of the 1,106 violations the reported medium run records across 276 instances, and at 50% of the 120 that occur where the grammar admits exactly one token, where a positional DP has no alternative to write.

C System optimisations

The repair layer sits on a frontier-constrained decoder, and the latencies we report depend on how that decoder avoids redundant work. Parser work is not repeated: one incremental parser is carried across denoising steps rather than re-parsing the prefix each time, and the search probes candidates by advancing and rolling back, so a parser is copied only for states that survive. Model work is not repeated either: resampling reuses the logits already in hand instead of a fresh forward pass, and committed tokens are batched: a clean step doubles the batch up to a cap of eight and any violation resets it to one, with the batch held at one until three quarters of the denoising steps have elapsed so that early positions each get a forward pass of context before they are fixed. Finally, mask construction for the next position overlaps the current forward pass, keeping the parser off the critical path.

D Implementation and hyperparameters

Every constant the decoder uses, in one place, so that the method is specified without reading it out of prose. Values are those of the reported runs.

	symbol	value	where it acts
<i>Backbone and schedule</i>			
denoising steps	T	128	
generation length		256 tokens	
block length		32	semi-autoregressive blocks
temperature		0.2	Gumbel noise on the logits
remasking		low-confidence	LLaDA’s reference sampler
seed		0	one seed, see §6
hardware		one A100-80GB per shard	
<i>Repair layer</i>			
candidate budget	k	100	per lattice node (§3.4)
span budget	L	48 positions	forward scan cap (§3.1)
DP position cap		64	never binds while $L < 64$
single-token retry depth		10	ranked logits at v (§3.1)
window policy		full span	one pass, not progressive widening
merge key	κ	$\text{bytes}(\mathcal{L}(q))$	§3.3, App. F
paths kept per key		1	widening it changes nothing measured
<i>Budgets and batching</i>			
hand-backs per instance		100	a run that exceeds it is reported as is
DP wall-clock per instance		240 s	never binds: scoring caps at 120 s
commit batch cap		8 tokens	
commit batch growth		$\times 2$ on a clean step	reset to 1 on any violation
batch growth gated until		$\frac{3}{4}T$	early positions get a pass each
<i>Parser</i>			
library		LLGUIDANCE 1.7.0	App. F
<code>max_items_in_row</code>		20,000	default aborts on deep nesting
<code>step_max_items</code>		600,000	
stop tokens		EOS 126081, EOT 126348	excluded from $\mathcal{C}$
mask token		126336	
vocabulary	$ V $	126,349	from the checkpoint’s tokenizer

Two behaviours the table cannot carry. First, a rejected single-token retry does not count against the hand-back budget: only a commit or a fall-through to the joint search does, so a violation whose repair sits below rank 10 cannot exhaust the budget before the search is reached. Second, mask construction for the next position is issued asynchronously and joined after the current forward pass, and any state change at the frontier discards the pending mask rather than using a stale one.

E What the search guarantees

The properties the repair layer relies on are stated separately here, since they hold under different conditions and two of them are approximate.

1. Local soundness, unconditional. Every token the search places is one the parser consumed from the state that precedes it: candidates are drawn from $\mathcal{L}(q)$ and the backtrace follows predecessor links rather than re-deriving a path from keys. No setting of k , L or the merge width can make the search emit a span the grammar rejects.

2. Termination, and an answer where one exists. The pass halts after at most one layer per position of the span. $\mathcal{C}$ is empty only where the grammar admits nothing but a stop token, which §3.2 disposes of before the search runs, and no live state we measured had an empty legal set (0 of 6,657, Appendix A). Where a later layer empties, the pass returns the prefix it has and the remaining positions keep the tokens the model wrote, which Algorithm 1 treats as the span finishing rather than failing.

3. Optimality, under exact state identity. Were κ injective on reachable parser states, the pass would return the maximiser of $\sum_p \log p_\theta$ over the candidate lattice, by the usual Viterbi argument. Two approximations separate that from the maximiser over all valid assignments of the span: the candidate budget k , and κ itself.

4. The merging approximation, measured and not bounded. κ is not injective; Appendix F gives a schema on which four states owing different numbers of array elements share a mask. Merging can therefore discard a higher-scoring path, though never a valid one, by (1). Widening the merge to four paths per key changes the returned assignment at 0.33% of solves. That is a lower bound on the loss and not a bound on it: the only comparison against a genuinely unmerged search is on a lattice small enough to enumerate, where the merged search returned the lattice optimum in all 7,836 solves.

5. Complexity, output-sensitive. Each layer holds at most $|S|$ states and expands at most k candidates from each, so the pass costs $O(\text{span} \cdot |S| \cdot k)$ parser operations and $O(|S|)$ live parsers. $|S|$ is the number of distinct masks the grammar can present at one position; nothing bounds it a priori, and it stayed in the single digits in every run reported here. Span length is not an exponent for any $|S|$.

F The parser interface

The paper treats the incremental parser abstractly: a state q , a consumable set $\mathcal{L}(q)$, and the ability to advance, copy, and undo. This appendix records the concrete interface those abstractions are built on, so that the method can be reimplemented against a different parser.

Construction. We use LLGUIDANCE 1.7.0 [Microsoft, 2025]. A tokenizer is built from the checkpoint’s own `tokenizer.json`, so token boundaries in the grammar and in the model agree exactly. A JSON Schema is compiled with `grammar_from_json_schema` and a SMILES grammar with `grammar_from_lark`; compilation is checked before use, and a schema that fails to compile is excluded from every arm rather than scored as a failure for one (§4 gives the counts). The parser limits are raised above their defaults, at which deeply nested schemas abort mid-parse (Appendix D).

The five operations.

paper	LLGUIDANCE	note
$\mathcal{L}(q)$	<code>compute_logit_bias()</code>	byte mask over V ; 0.025 ms mean
advance	<code>try_consume_tokens([t])</code>	returns tokens consumed; 0 means rejected
copy a state	<code>deep_copy()</code>	lattice nodes hold these
undo	<code>rollback(n)</code>	probe without copying
may stop	<code>is_accepting()</code>	state is final, not that it must stop

Two details the method depends on. First, EOS is marked legal by the bias at accepting states but is not a consumable grammar symbol: `try_consume_tokens([EOS])` returns 0 there. It is therefore excluded from the candidate set of Eq. 1: offering it wastes a slot, and where it is the only token the bias marks legal it would empty the set and imitate a dead end. A termination rule written in terms of consumption would therefore never fire, which is why `only_stop_remains` (§3.2) is defined on the bias, meaning that nothing but stop tokens remains legal, rather than on what the parser will consume. Second, the DP merges lattice paths on `bytes(compute_logit_bias())` as the state key. The key is a value the implementation already computes to build the candidate set, so merging costs no additional parser calls, but it identifies a state by what that state admits next rather than by the parser’s configuration, and the two are not the same thing. An array with `"minItems": 5` gives the separating case, and it is easy to check. Feeding `{"xs": [1, {"xs": [1, 2, {"xs": [1, 2, 3 and {"xs": [1, 2, 3, 4` to the parser leaves four states that owe four, three, two and one more element; all four return byte-identical masks, since all four must write a comma next and none may close the bracket. The mask separates them only once the obligation is discharged: after a fifth element it admits `]` and differs from all of them. Merging on κ therefore collapses states that still differ, keeping one and discarding the rest. §3.3 states what that collapse can and cannot cost.

Collisions are rarer than the definition suggests, because the mask is over whole tokens of a byte-level BPE rather than over grammar symbols: a token spanning several bytes makes the mask reflect what the grammar permits several bytes ahead, so states whose futures diverge within that reach are already separated by their masks. A collision requires the divergence to lie beyond it, as it does above. The key can be refined at no cost, since conjoining `is_accepting()` or the bracket depth the span-end scan already maintains separates states without an additional parser call, and we report the unrefined key because it is what produced the measurements below.

How much the abstraction costs. We measure it by widening the merge rather than by reasoning about it. The DP takes a width B : the best B paths per key survive instead of the best one, so $B=1$ is the reported configuration and $B \rightarrow \infty$ is no merging at all. Widening strictly enlarges the searched set, so a width at which the answer stops moving is a width at which merging costs nothing. We replay the generated outputs of the `medium` run through the parser and, at every structural state reached, outside string literals and with a legal set under 4,000, solve the same span twice, varying only B , with the model’s ranking replaced by random score vectors so that the merge has to survive many orderings of the same candidate sets rather than one. Over 13,513 solves at 4,529 states of 447 instances, $B=2$ changed the assignment at 0.19% and $B=4$ at 0.33%, the second set containing the first; at every one of those the wider search scored strictly higher, by a median of 4.6 nats and at most 21.2. The merge therefore does discard a likelier path, on roughly one solve in three hundred. It discards a *valid* path never, which is the property the layer depends on. Where the lattice is small enough to run with no merging at all ($k=6$ over four positions, 7,836 solves in which the path cap never bound), the merged search returned the lattice optimum every time, so what the key loses at $k=100$ is reachable only through candidate sets wider than a handful. For scale, the unmerged arm carried a median of 1,869 live paths where $B=1$ carried 5.

566 G Prevention on the subset where prevention is available

The introduction sets repair against prevention, and the published prevention methods are the ones we cannot run (§4). Running them would mean compiling each JSON Schema to a finite automaton first, which the recursive schemas in these splits do not admit at all and which, on the ones that do, would measure that translation as much as the decoding algorithm; their guarantee is therefore a property we report rather than a number we measure. So that the contrast is measured rather than asserted, we implement the prevention side ourselves, on the same backbone, the same prompts and the same schedule. At each denoising step it conditions the model’s mean-field prediction on the whole constraint and decodes from the exact constrained posterior

$$p(x^0 \mid x^t, G) \propto \left(\prod_i p_{\theta}(x_i^0 \mid x^t) \right) \cdot \mathbf{1}[x^0 \in \mathcal{L}(G)],$$

which a finite automaton makes tractable: the automaton is a chain over positions, multiplying a chain by a fully factorised measure leaves a chain, and exact inference is then forward-backward. Violations cannot occur, so none of the repair machinery is reachable. We report both decoders, because the distinction matters here: *viterbi* takes the mode of that posterior and *marginal* takes per-position marginals.

580 The constraint is compiled to an automaton under a nesting bound and an edge budget of 20 million,
581 which is where a recursive schema is lost, and the arms are scored on 101 instances rather than on the
582 full split because an A100 minute buys about one instance here. Table 5 reports what the guarantee
583 costs.

Table 5: Prevention against repair on the first 128 instances of JSONSchemaBench medium in identifier order, less 27 on which automaton construction exhausted a 32 GiB container before any decoder ran, leaving 101; same backbone, prompts and schedule throughout. Rates are over all 101, so a schema an arm’s toolchain cannot accept counts against it: the guarantee the prevention arms carry is over the documents they produce, not over the workload. Constraint time is each arm’s own instrumentation and the two harnesses measure it differently, so it is compared as an order of magnitude rather than as a ratio.

Method	constraint built	schema@1 ↑	content ↑	median leaves ↑	mean s ↓
FA, marginal (mass)	66 (65.3%)	59.4	8.9	1	77.5
FA, viterbi (mode)	67 (66.3%)	58.4	18.8	1	57.5
no-DP (ours)	92 (91.1%)	83.2	67.3	13	13.4
DPGrammar	92 (91.1%)	85.1	73.3	13	16.8

584 Prevention pays on all three axes at once. It reaches a quarter fewer schemas: an automaton is built
585 for 66% of the sample against the incremental parser’s 91%. The gap is not a tuning choice but the
586 size of the object being asked for. Of the 34 schemas the construction does not deliver, 16 exceed the
587 edge budget, at a median of 61.7 million edges and one schema at 2.0 billion against a budget of 20
588 million, and 16 fail construction outright; the remaining two decode to nothing usable. The recursion
589 a JSON Schema expresses in a line is what a finite automaton has to pay for in states, and on a
590 third of this workload it cannot. It is slower by a factor of three to five, and where the repair layer’s
591 constraint machinery is under a second of a 16.8 s instance, the prevention arms spend 90.5% and
592 92.4% of wall clock inside inference over the automaton. And on the documents it does produce it is
593 degenerate: the median prevention output carries a single scalar leaf against DPGrammar’s thirteen,
594 and only 8.9% of the mass decoder’s and 18.8% of the mode decoder’s outputs clear the five-leaf
595 floor against 73.3% of ours. This is the failure the floor was introduced to see (§4): the smallest
596 document a schema admits is both valid and cheap, so a decoder that optimises probability under a
597 hard constraint is drawn to it, and closing a structure costs one token where content pays entropy at
598 every position.

599 One number here is worth reading against the guarantee rather than past it. Validity is 88–91% over
600 the instances the prevention arms actually run, not 100%, and the shortfall is not the decoder. Every
601 output is a string the automaton accepts; the automaton is simply not the schema. Compiling a JSON
602 Schema to a finite automaton under a nesting bound is lossy in the permissive direction, and of the
603 eight invalid outputs among the 67 the mode decoder produced, four place a string where the schema
604 requires an object, one exceeds a numeric maximum, one violates a pattern, one satisfies no branch
605 of an anyOf, and one is not JSON at all. None is truncated. Prevention is exact with respect to the
606 automaton it is given, which makes its guarantee only as good as that compilation, and this is the
607 fourth cost the guarantee carries alongside coverage, latency and content.

608 H Comparing on shared populations

609 CD-CFG [Mündler et al., 2025a] is run from the authors’ released code, which builds its constraint
610 with RUSTFORMLANG, the formal-language library shipped in the same repository, rather than with
611 LLGUIDANCE [Microsoft, 2025]. Counting the schemas that front-end rejects as decoding failures
612 would measure a regex parser rather than a decoding algorithm, so Table 6 reports coverage separately
613 and scores validity only on the 351 medium schemas both toolchains compile. Coverage differs
614 widely: RUSTFORMLANG rejects 203 of 586, mostly over regular-expression quantifiers that Python’s
615 re accepts, where LLGUIDANCE accepts 511.

616 On the shared 351, CD-CFG reaches 66.1% with 19 timeouts against DPGrammar’s 97.4% with none,
617 and every arm scores higher there than on its own full set, so the 160 schemas only LLGUIDANCE
618 compiles are the harder ones.

Table 6: Grammar coverage is a property of the compiler, not of the decoder. Validity is measured only where both toolchains compile, so the two effects are not confounded.

Method	Toolchain	Schemas compiled	Validity on the shared 351
CD-CFG	RUSTFORMLANG	383/586 (65.4%)	66.1%
LAVE	LLGUIDANCE	511/586 (87.2%)	82.3%
No-DP baseline	LLGUIDANCE	511/586 (87.2%)	92.6%
DPGrammar	LLGUIDANCE	511/586 (87.2%)	97.4%

CD-CFG is scored with the `autocomplete_valid` post-processing its own released code applies after decoding, which closes a document the decoder left open. That step is worth about twenty points and we report the arm with it: on `medium` it takes CD-CFG from 44.1% to 64.8%, on `easy` from 56.3% to 79.1%, on `hard` from 11.4% to 32.2%, and on the shared 351 from 45.0% to 66.1%. Without it the arm falls below the unconstrained decoder at every tier; with it, it does not. Our own arms use no comparable step: a DPGrammar document is closed by the grammar or not at all. LAVE needs the same treatment for a different reason: it runs on the same 511 we do but reaches the cap on 60, and judging it only where it finishes raises it from 78.3% to 88.7% against our 97.3% on that same 451. Neither adjustment changes the ordering.

I Other decoding schedules

The repair layer fires only after a parser rejection, so it places no requirement on the order in which positions are revealed. That is a claim about applicability, and Table 7 separates it from a claim about the size of the gain. Both arms are rerun end to end under each schedule, so every row is an internally controlled comparison.

Table 7: The same two arms under three decoding schedules, on the 511 `medium` instances LLGUIDANCE compiles. *sched. end* is the share of instances whose generation stopped because the $T=128$ denoising schedule ran out rather than because the grammar permitted an end.

Schedule	Arm	schema@1	content	rs/inst	violations	DP dead ends	sched. end
low-confidence, block 32	no-DP	91.8	74.6	33.2	—	—	—
	DPG	96.7	84.9	1.2	1,106	3/315	31%
random unmasking order	no-DP	91.6	59.1	63.3	—	—	—
	DPG	92.0	76.3	3.1	2,576	1/733	46%
full parallel, block 256	no-DP	77.3	63.8	30.0	—	—	—
	DPG	92.4	76.9	1.0	872	0/320	39%

Committing the whole sequence in one block rather than in blocks of 32 puts more of it beyond the frontier before the parser reads it, and the validity margin widens from +4.9 to +15.1 pp: the unrepaired baseline falls to 77.3% while the repaired arm holds 92.4%. This is the difficulty gradient of Table 1 arriving along a second axis.

Random order costs both arms, and for a reason the table makes visible. It produces $2.3\times$ the violations of the reported schedule, each of which costs denoising steps to repair, so 46% of instances exhaust the 128-step schedule before the document closes against 31% here. Those truncations are 98% of what DPGrammar fails on under that schedule. The search itself does not fail: it dead-ends once in 733 calls, against three in 315 under the reported schedule. What the fixed step budget costs is validity, not repair.

The content margin survives every schedule and is largest where the schedule is hardest, at +17.2 and +13.1 pp against +10.4 in the reported configuration. The validity margin does not: it ranges from +0.4 to +15.1, tracking how much schedule is left once repair has been paid for. Raising T is the obvious remedy and one we do not test here.

647 J Where the extra wall time goes

648 DPGrammar is 3.3 s slower than the unrepaired baseline on medium. Figure 3 attributes that difference
 649 to work rather than to a time breakdown, because the instrumented counters cover only about a
 650 quarter of wall clock in either arm.




Figure 3: Averages per instance on JSONSchemaBench medium, split by whether the repair layer ever fired. (a) The mean difference is concentrated in the 173 instances where it did; on the other 338 the layer is slightly faster. (b) On those instances it emits 71% more tokens and needs proportionally more forward passes, while remasking 28 $\times$ less. (c) Per unit the two arms are close, and the constraint machinery costs the same in both.

651 The bottleneck is not the search. The constraint machinery costs 909 ms against the baseline’s 920,
 652 and grammar checks 22.5 ms against 11.2. What changes is how much gets written: the baseline
 653 stalls, remasking 79 times and stopping at 133 tokens, while DPGrammar repairs and continues to
 654 227. That takes 1.5 $\times$ the forward passes, each 1.2 $\times$ slower on the longer sequence, which is the
 655 whole of the difference.

656 K Why validity alone is not enough

657 Schema validity can be blind to a difference that matters, which is why §4 pairs it with a content floor
 658 rather than reporting it alone. On o78738 the schema asks for a course catalogue entry. Both arms
 659 validate, so schema@1 scores them identically; what separates them is that the unrepaired baseline
 660 stalls inside prereqs, emits two hundred lines of whitespace, and closes the array empty:

```
no-DP: valid, 4 populated leaves
[ { "code": "MATH 1001", "name": "Calculus I", "credits": "4",
  661   "description": "An introduction to calculus",
   "requirements": { "prereqs": [
                                <200 lines of whitespace>
                              ], "coreqs": [] } } ]
```

662 DPGrammar reaches the same violation at position 120, solves a three-position span in which the
 663 parser admits fewer than two tokens on average, and the decoder continues through the rest of the
 664 record:

```
DPGrammar: valid, 12 populated leaves
[ { "code": "MATH 1001", "name": "Calculus I", "credits": "4",
  665   "description": "An introduction to calculus",
   "requirements": {
     "prereqs": [ ["MATH 0001", "MATH 0002"] ],
     "coreqs": [ ["MATH 1001", "MATH 1002"] ],
     "lectureHours": "3", "tutorialHours": "-", "labHours": "2",
     "note": "Recommended for first-year students" } } ]
```

666 Neither output is wrong about the course; the baseline simply stops saying anything, and validity
 667 cannot see the difference. Only the content floor separates them, and it does so by 8 populated leaves.

668 Neither instance is exceptional: across medium, 18.8% of the baseline’s valid outputs carry fewer
 669 than five scalar leaves, against 12.1% of ours.